

MANSI DHRUV

Lead Data Engineer · Solution Architect

Maharashtra, India · mansi.p.dhruv@gmail.com · [linkedin.com/in/mansidhruv](https://www.linkedin.com/in/mansidhruv) · github.com/MsMansiDhruv · portfolio-mansi-eight.vercel.app

SUMMARY

Lead Data Engineer with 7+ years building production data platforms, ETL/ELT, and cloud architectures on AWS and Databricks. Most of my work starts as a migration or cost problem and becomes an architecture problem: layered lakehouses, workload-specific serving, and ML paths that data scientists can trust in production. I lead delivery, make trade-offs explicit, and mentor engineers so the system stays operable after it ships.

SKILLS

Data platforms	Spark · PySpark · Scala · SQL · Databricks · Delta Live Tables · Airflow · Kafka · Kinesis
AWS	S3 · Glue · Redshift · Lambda · DynamoDB · ECS · EC2 · API Gateway · DMS · Athena
Engineering	Python · Terraform · CloudFormation · Docker · CI/CD · FastAPI · GraphQL · Bitbucket Pipelines
Analytics / ML	Power BI · Tableau · MLflow · production ML pipeline integration · web extraction (Scrapy, Selenium)

EXPERIENCE

Lead Data Engineer

Oct 2021 – Present

SG Analytics · Maharashtra, India

Progressed Data Engineer → Senior Data Engineer → Lead Data Engineer (2025). Own architecture, delivery, and mentoring across confidential enterprise engagements.

- Lead a cross-functional data team: architecture decisions, delivery standards, and mentoring of engineers and interns in Python, production practices, and automation.
- Modernized legacy shell ETL into a layered AWS lakehouse (S3 Bronze/Silver/Gold, Glue/PySpark, Redshift, Power BI) for a confidential asset-management client — centralized analytics and self-service reporting.
- Separated high-frequency application serving from warehouse analytics after a Redshift cost signal; benchmarked Redshift, Aurora, DynamoDB, and S3 Tables and chose engines by workload instead of one universal store.
- Productionized an ML-driven allocation path (Databricks, Delta, MLflow, ECR, GraphQL) for a confidential fintech engagement — engineering the serving pipeline around Data Science-owned models.
- Built batch and streaming pipelines (Kafka, Kinesis, Delta Live Tables, Spark/Scala) on AWS; documented ~40% end-to-end pipeline-time reduction and ~30% infrastructure-cost reduction versus the prior Databricks setup.
- Automated a web-intelligence pipeline (Scrapy, Selenium, Lambda, Terraform, CI/CD) from discovery through deduplication, classification, and reporting; reduced manual monitoring workload by ~50%.
- Delivered Glue/Redshift ETL and AWS infrastructure for an enterprise life-sciences client, including dashboard-facing transforms and ML-model data prep; automated processing and cut data-processing time by ~45%.
- Diagnosed slow SQL inserts with structured RCA and query redesign, improving time efficiency by ~85% on the affected paths.

Software Engineer

Aug 2019 – Oct 2021

CypherTree Technologies · Maharashtra, India

- Built backend systems for a self-service social-analytics platform (Django, React, Redshift, Celery); shipped REST APIs and reduced API load time by 25–30%.
- Implemented payment-gateway integrations and GraphQL/JWT services for an e-commerce platform (Python, Django, Elasticsearch, Docker).

Software Engineering Intern

Dec 2018 – Jul 2019

Bemrr · Gujarat Police · Gujarat, India

- Built production iOS and web applications (Django, Python, JavaScript, Ionic).

EDUCATION

(Integrated) Master of Computer Applications · Ahmedabad University

2014 – 2019

Score: 3.94 / 4.33

CERTIFICATIONS & AWARDS

Certifications: AWS Certified Cloud Practitioner (Aug 2022) · Building Your First ETL Pipeline Using Azure Databricks (2022) · Big Data Analytics using Spark (2021)

Awards: Value-able Award, SG Analytics (2024) — foundational delivery that led to new client business · GEM Award, SG Analytics (2023) — exceptional business value